

Part 2 Submission: Role-Governed Chatbot System

1. Role-Based Chatbot Flow Structure

The chatbot architecture leverages an ****Isolated Intelligence Silo**** model. Instead of utilizing a single monolithic bot traversing varied conditionals, each authenticated role is dynamically mapped to a specialized agent instance. These instances are governed by distinct personalities and limited functional domains.

Agent Class	Target Role	Persona Profile	Functional Expertise
BuyerBot	BUYER	Residential Sales Expert	Floorplans, pricing, EMI options, booking, availability
PartnerBot	CHANNEL_PARTNER	Partner Relations Manager	Commission structures, referral policies, payouts
VisitBot	SITE_VISIT	Site Visit Coordinator	Visit scheduling, logistics, directions
EnquiryBot	ENQUIRY	General Enquiry Specialist	Developer legacy, high-level project overview

2. Routing logic explanation

The system utilizes an initial ****Verified Identity Check**** to establish a secure and dedicated conversational domain:

- **Identity Acquisition:** The user supplies a unique identifier, primarily a phone number.
- **Database Lookup:** The routing engine queries the verified classification data (`classified_leads_output.csv` from Part 1) to retrieve the official Assigned_Role mapped to that user.

- **Instantiation Locking:** Based on the result, a factory pattern initializes the corresponding specialized Agent object (e.g., BuyerBot). The session state acts as a hard boundary, effectively "locking" the user's entire interaction workflow to that strictly defined class. A user identified as a "Buyer" cannot route their query through the "Partner" module regardless of their phrasing.

3. Knowledge isolation strategy

Information security and compartmentalization are enforced via a stringent **Triple-Layer Strategy**:

- **Layer 1: Contextual Memory Constraints:** Each Agent class contains its own dedicated knowledge dictionary upon initialization. A `BuyerBot` simply does not have commission data injected into its context array.
- **Layer 2: Prompts and Persona Scoping:** The core LLM system prompt explicitly mandates the Bot to respond **only** based on its localized knowledge string, ensuring the AI model acknowledges its boundary limitations.
- **Layer 3: Pre-Response Moderation (Guardrails):** All incoming user queries are algorithmically scrubbed against a "Banned Keyword" array unique to the bot's class **before** LLM inferencing. If a keyword breach occurs (e.g., a Site Visit lead asking about financial payouts), a declarative static refusal is returned, preventing the AI reasoning engine from even parsing the out-of-bounds text.

4. Risk mitigation strategy

Maintaining data integrity and user authorization requires the following proactive security principles:

- **Preventing Cross-Role Data Leakage:** By employing pre-response algorithmic guardrails, sensitive topics are intercepted prior to reaching the language processing layer. This makes unauthorized data hallucination technically implausible within the context window.

- **Handling Incorrect Role Data:** Discrepancies between expected roles and assigned database roles are mediated by a direct ****Reporting Loop****. If the system denies a user access due to incorrect pre-classification, they utilize the "Report Role Mismatch" function. This triggers a priority flag (`ROLE_MISMATCH_REPORTED`) in the admin audit log for rapid remediation.
- **Controlling Hallucinations:** Agent output generation is firmly contextualized within restricted local memory arrays. The system rejects extraneous conversational threads, reducing the risk of generating inaccurate business logic.
- **Compliance Auditability:** Every interaction node—authorized queries and blocked violations—is asynchronously written to `interaction_audit.json` with associated metadata (time, role, phone, policy flag). This creates a persistent and verifiable telemetry trail for compliance supervisors.

4. Architecture diagram



